


Hash Table

Introduction

- **Dictionary**: a dynamic set that supports INSERT, SEARCH, and DELETE
- **Hash table**: an effective data structure for dictionaries
 - $O(n)$ time for search in worst case
 - $O(1)$ expected time for search

Direct-Address Tables


- An Array is an example of a direct-address table
 - Given a key k , the corresponding is stored in $T[k]$
 - Assume no two elements have the same key
 - Suitable when the universe U of keys is reasonably small
 - $U = \{0, 1, 2, \dots, m-1\} \rightarrow T[0..m-1] \rightarrow \Theta(m)$
 - What if the no. of keys actually occurred is small
- Operations of direct-address tables: all $\Theta(1)$ time
 - $\text{DIRECT-ADDRESS-SEARCH}(T, k)$: return $T[k]$
 - $\text{DIRECT-ADDRESS-INSERT}(T, x)$: $T[\text{key}[x]] \leftarrow x$
 - $\text{DIRECT-ADDRESS-DELETE}(T, x)$: $T[\text{key}[x]] \leftarrow \text{NIL}$




If the set contains no element with key k , $T[k] = \text{NIL}$

Implementing a dynamic set by a direct-address table T . Each key in the universe $U = \{0, 1, \dots, 9\}$ corresponds to an index in the table. The set $K = \{2, 3, 5, 8\}$ of actual keys determines the slots in the table that contain pointers to elements. The other slots, heavily shaded, contain NIL.

Alternative Illustration of Direct-Address Tables



Use separate arrays for key and satellite data if the programming language does not support objects (parallel array)



If the programming language can support object-element arrays


Hash Tables

Overview

- Difficulties with direct addressing
 - If U is large, storing a table T of size $|U|$ may be impractical
 - The set K of keys actually stored may be small relative to U
- When the set K of keys stored in a dictionary is much smaller than U of all possible keys, a hash table requires much less storage than a direct-address table
 - $\Theta(|K|)$ storage
 - $O(1)$ average time to search for an element
 - $\Theta(1)$ worst case to search for an element in direct-address table

Hash Table


- Direct addressing: an element with key $k \rightarrow T[k]$
- Hash table: an element with key $k \rightarrow T[h(k)]$
 - $h(k)$: hash function; usually can be computed in $O(1)$ time
 - $h: U \rightarrow \{0, 1, \dots, m-1\}$
 - An element with key k hashes to slot $h(k)$
 - $h(k)$ is the hash function of key k
 - Instead of $|U|$ values, we need to handle only m values
 - Collision may occur
 - Two different keys hash to the same slot ($h(k_1) = h(k_2) = x$)
 - Eliminating all collisions is impossible
 - But a well-designed “random”-looking hash table can minimize collisions



Using a hash function h to map keys to hash-table slots. Keys k_2 and k_5 map to the same slot, so they collide.

Collision Resolution by Chaining

- Idea
 - Put all the elements that hash to the same slot in a linked list
 - Slot j contains a pointer to the head of the list of all stored elements that hash to j
- Operations
 - CHAINED-HASH-INSERT(T, x): $O(1)$ (no check for duplication)
 - Insert x at the head of list $T[h(key[x])]$
 - CHAINED-HASH-SEARCH(T, k): proportional to the list length
 - Search for an element with key k in list $T[h(k)]$
 - CHAINED-HASH-DELETE(T, x): $O(1)$ for doubly linked list
 - Delete x from the list $T[h(key[x])]$



另一種方式:open addressing

Collision resolution by chaining. Each hash-table slot $T[j]$ contains a linked list of all the keys whose hash value is j . For example, $h(k_1) = h(k_4)$ and $h(k_5) = h(k_2) = h(k_7)$.

Analysis of Hashing with Chaining for Searching

- Load factor $\alpha = n \text{ elements} / m \text{ slots}$ (負載因子)
 - Average number of elements stored in a chain
- Worst case: $\Theta(n)$ + time to compute the hash function
 - all n keys hash into the same slot
- Average performance
 - How well the hash function h distributes the set of keys to be stored among the m slots, on the average
 - Assume simple uniform hashing
 - Any given element is equally likely to hash into any of the m slots, independently of where any other elements has hashed to
 - The time required for a successful or unsuccessful search is $\Theta(1 + \alpha)$
 - $n = O(m) \rightarrow \alpha = n/m = O(m)/m = O(1)$

Hash Functions

Overview

- Interpreting keys as natural numbers $N=\{0, 1, 2, \dots\}$
 - If the keys are not natural number \rightarrow convert
For example “pt”: $p=112, t=116$, using radix-128 integer \rightarrow pt becomes $112*128+116=14452$
 - What makes a good hash function?
 - Each key is equally likely to hash to any of the m slots, independently of where any other key has hashed to
 - Rarely knows the probability distribution
 - Keys may not be drawn independently
 - Hash by division: heuristic
 - Hash by multiplication: heuristic
 - Universal hashing: randomization to provide provably good performance
-

Hash By Division

- $h(k) = k \bmod m$
 - $m = 12$ and $k=100 \rightarrow h(k) = 4$
 - Must avoid certain values of $m \rightarrow$ rely on characteristics of k values
 - m should not be a power of 2 $\rightarrow h(k)$ is the lowest-order bits of k
 - Good values for m are primes not too close to an 2^p
 - For example: $n=2000$ character string, we don't mind examining 3 elements in unsuccessful search \rightarrow allocate a hash table of size 701
701 is prime near $2000/3$, not near any power of 2
 $h(k)=k \bmod 701$

Hash By Multiplication

- $h(k) = \lfloor m * (k * A \bmod 1) \rfloor$
 - multiply the key k by a $A \in (0, 1)$ and extract the fractional part of kA
 - multiply the fractional part of kA by m and take the floor of the result
- The value of m is not critical
 - Typically, $m = 2^p$ (for easy implementation on computers)
 - (Knuth) $A \approx (\sqrt{5} - 1) / 2 = 0.6180339887..$
 - For example: $k=123456, m=10000, a=0.618$
 $H(k) = \text{floor}(10000 * (123456 * 0.618... \bmod 1))$
 $= \text{floor}(10000 * (76300.004151... \bmod 1))$
 $= \text{floor}(10000 * 0.\underline{0041}151.....) = 41.$

Universal Hashing

- Idea:
 - Choose the hash function randomly in a way that is independent of the keys that are actually going to be stored
 - Select the hash function at random from a carefully designed class of functions at the beginning of execution
 - The algorithm can behave differently on each execution, even for the same input

Universal Hashing (Cont.)

- Let H be a finite collection of hash functions that map a given universe U of keys into the range $\{0, 1, \dots, m-1\}$.
- H is **universal** if for each pair of distinct keys $k, l \in U$, the number of hash functions $h \in H$ for which $h(k) = h(l)$ is at most $|H|/m$
 - Collision chance = $1/m$
- Theorem 11.3.
 - Suppose that a hash function h is chosen from a universal collection of hash function and is used to hash n keys into a table T of size m , using chaining to resolve collisions. If key k is not in the table, then the expected length $E[n_{h(k)}]$ of the list that key k hashes to is at most α . If key k is in the table, then the expected length $E[n_{h(k)}]$ of the list containing key k is at most $1 + \alpha$

Designing A Universal Class of Hash Functions

- Steps to design a universal class of hash functions
 - Choose a prime number p , so that $k \in [0, p-1]$ and $p > m$
 - $Z_p = \{0, 1, \dots, p-1\}$
 - $Z_p^* = \{1, 2, \dots, p-1\}$
 - $h_{a,b}(k) = ((ak+b) \bmod p) \bmod m$, for any $a \in Z_p^*$ and $b \in Z_p$
- The family of all such hash functions is
 - $H_{p,m} = \{h_{a,b} : a \in Z_p^* \text{ and } b \in Z_p\}$
 - Total: $p(p-1)$ hash functions in $H_{p,m}$

Other hash function(1)

- Mid-square: 先將數值平方在取中間部分位元

$$(39)^2 = (100111)^2 = (10111110001)_2$$

$$F(39) = (11110)_2 = (29)_{10}$$

- Digital analysis: 位數分析, 對資料每一位數加以分析, 剔除不均勻分布之位數, 剩下位數作為hash位址


適合static data 且位數相同

For example

x	F(x)
1510327	5107
1857384	8574
2621439	6219
1796333	7963
1038420	0380
7142481	1421
7203326	2036
2385425	3855

Other hash function(2)

- Folding(摺疊法):
- (1)shift folding
- (2)folding at boundaries



Open Addressing

Overview

- All elements are stored in the hash table itself
 - Each table entry contains either an set element or NIL
 - Search: systematically examine table slots until the desired element is found or it is clear that the element is not in the table
 - No lists and no elements are stored outside the table
 - Load factor $\alpha \leq 1$
 - At most m elements can be stored in the hash table
 - The extra memory freed by not storing pointers provides the hash table with a larger number of slots for the same amount of memory, potentially yielding fewer collisions and faster retrieval

Probe

- **Insertion**: successively examine (*probe*) the hash table until we find an empty slot in which to put the key
 - The sequence of positions probed relies on the key being inserted
 - *Probe sequence* for every key k must be a permutation of $\langle 0, 1, \dots, m-1 \rangle$
 - Extended hash function $\rightarrow h: U * \{0, 1, \dots, m-1\} \rightarrow \{0, 1, \dots, m-1\}$
 - $\langle h(k, 0), h(k, 1), \dots, h(k, m-1) \rangle$
- Assume uniform hashing
 - Each key is equally likely to have any of the $m!$ permutations of $\langle 0, 1, \dots, m-1 \rangle$ as its probe sequence

INSERT and SEARCH in Open Addressing

HASH-INSERT(T, k)

```
1   $i \leftarrow 0$ 
2  repeat  $j \leftarrow h(k, i)$ 
3          if  $T[j] = \text{NIL}$ 
4              then  $T[j] \leftarrow k$ 
5                  return  $j$ 
6              else  $i \leftarrow i + 1$ 
7  until  $i = m$ 
8  error "hash table overflow"
```

HASH-SEARCH(T, k)

```
1   $i \leftarrow 0$ 
2  repeat  $j \leftarrow h(k, i)$ 
3          if  $T[j] = k$ 
4              then return  $j$ 
5               $i \leftarrow i + 1$ 
6  until  $T[j] = \text{NIL}$  or  $i = m$ 
7  return NIL
```

Assume keys are not
deleted from the table

Example

$$h'(k) = k \bmod 5$$

$$h(k,i) = (h'(k) + i) \bmod 5$$

$$\text{For } k=9 \rightarrow h(k,i) = \langle 4, 0, 1, 2, 3 \rangle$$

$$\text{For } k=3 \rightarrow h(k,i) = \langle 3, 4, 0, 1, 2 \rangle$$

0	14
1	
2	2
3	
4	9

HASH-INSERT(T,2)

HASH-INSERT(T,9)

HASH-INSERT(T,14)

0	14
1	
2	2
3	
4	DELETED

HASH-DELETE(T, 9)

HASH-SEARCH(T, 14)

We have to mark a deleted slot,
instead of letting it be NIL

DELETE in Open Addressing

- Difficult !!
 - Cannot simply mark the deleted slot i as empty by storing NIL
 - Impossible to retrieve any key k during whose insertion we had probed slot i and found it occupied
 - Solution: mark the slot by storing in it a special value DELETED
 - Modify HASH-INSERT: treat a DELETED slot as a NIL slot
 - No modification of HASH-SEARCH is needed
 - Search times are no longer dependent on α
 - Therefore, chaining is more commonly selected as a collision resolution technique when keys must be deleted

Linear Probing

- $h(k,i) = (h'(k) + i) \bmod m$ ($i = 0, 1, \dots, m-1$)
 - $h': U \rightarrow \{0, 1, \dots, m-1\}$ is called an auxiliary hash function
 - $T[h'(k)] \rightarrow T[h'(k)+1] \rightarrow \dots \rightarrow T[m-1] \rightarrow T[0] \rightarrow \dots \rightarrow T[h'(k)-1]$
 - Only m distinct probe sequences
 - Easy to implement
 - Problem: Primary clustering
 - Long runs of occupied slots tend to get longer, and the average search time increases
 - Secondary clustering
 - If two keys have the same initial probe position, then their probe sequences are the same
 - Example: $f(x)=x$, table size=19(0..18),
data={1.0.5.1.18.3.8.9.14.7.5.5.1.13.12.5}

Quadratic Probing

- $h(k,i) = (h'(k) + c_1i + c_2i^2) \bmod m$ ($i = 0, 1, \dots, m-1$)
 - h' : auxiliary hash function; $c_1, c_2 \neq 0$ (auxiliary constants)
 - The values of c_1, c_2, m are constrained (See Problem 11-3)
 - Work much better than linear probing (alleviate primary clustering)
 - Secondary clustering
 - If two keys have the same initial probe position, then their probe sequences are the same
 - Only m distinct probe sequences

Double Hashing: the best

- $h(k,i) = (h_1(k) + i \cdot h_2(k)) \bmod m$ ($i = 0, 1, \dots, m-1$)
 - h_1 and h_2 are auxiliary hash functions
 - Depends in two ways upon the key i , since the initial probe position, the offset, or both, may vary $\rightarrow \Theta(m^2)$ distinct probing sequences
 - The value of $h_2(k)$ must be relative prime to the hash-table size m for the entire hash table to be searched (Exercise 11.4-3)
 - Let $m = 2^p$ and design h_2 so that it always produces an odd number
 - Let m be prime and design h_2 so that it always produces a positive integer less than m
 - $h_1(k) = k \bmod 701$
 - $h_2(k) = 1 + (k \bmod 700)$

Example of Double Hashing

Data sequence: 69, 72, 50, 79, 98, 14

0	
1	79
2	
3	
4	69
5	98
6	
7	72
8	
9	14
10	
11	50
12	

$$h_1(14) = 1$$

$$h_2(14) = 4$$

$$H_1(98) = 98 \bmod 13 = 7$$

$$H_2(98) = 1 + (98 \bmod 11) = 11$$

$$H(98, 1) = (7 + 11) \bmod 13 = 5$$

$$h(k, i) = (h_1(k) + i * h_2(k)) \bmod m$$

$$h(14, 0) = 14 \bmod 13 = 1 \dots \dots \dots 1^{\text{st}} \text{ collision}$$

$$h(14, 1) = (1 + 1 * (1 + (14 \bmod 11))) \bmod 13 = 5 \dots \dots 2^{\text{nd}} \text{ collision}$$

$$h(14, 2) = (1 + 2 * (4)) \bmod 13 = 9$$

$$H_1(k) = k \bmod 13$$

$$H_2(k) = 1 + (k \bmod 11)$$

Insertion by double hashing. Here we have a hash table of size 13 with $h_1(k) = k \bmod 13$ and $h_2(k) = 1 + (k \bmod 11)$. Since $14 \equiv 1 \pmod{13}$ and $14 \equiv 3 \pmod{11}$, the key 14 is inserted into empty slot 9, after slots 1 and 5 are examined and found to be occupied.

Analysis of Open-Address Hashing

- Theorem 11.6
 - Given an open-address hash table with load factor $\alpha = n/m < 1$, the expected number of probes in an unsuccessful search is at most $1/(1-\alpha)$, assuming uniform hashing
 - If α is a constant \rightarrow an unsuccessful search runs in $O(1)$ time
 - Corollary 11.7
 - Inserting an element into an open-address hash table with load factor α requires at most $1/(1-\alpha)$ probes on average, assuming uniform hashing
 - Theorem 11.8
 - Given an open-address hash table with load factor $\alpha = n/m < 1$, the expected number of probes in a successful search is at most $1/\alpha * \ln(1/(1-\alpha))$, assuming uniform hashing and that each key in the table is equally likely to be searched for
 - If α is a constant \rightarrow an unsuccessful search runs in $O(1)$ time
-

Self-Study

- Proof of Theorems
- Section 11.5 Perfect Hashing
 - The worst-case number of memory accesses required to perform a search is $O(1)$

example

- A hashed table is constructed using the division hash algorithm function with 5 buckets (a bucket at most 4 records) . If the following key field values are to be placed in buckets: 3 ,5, 24, 22, 109, 10, 8, 6, 23, 28, 100, 103, 9, 39, 27, 0.
 - Identify the number of records in each bucket.
 - Which bucket overflows?

$$H(k)=k \bmod 5$$

•Bucket 3 overflow.

no
↓

0	5	10	100	0
1	6			
2	22	7	27	
3	3	8	23	28
4	24	109	9	29

103